

Math Nuggets for GEN AI

Retrieval-Augmented Generation: Mathematical Foundation of Knowledge-Enhanced LLMs

By Dr. Anish Roychowdhury, Plaksha University

September 23, 2025

Beyond Parametric Knowledge: The External Memory Challenge

Traditional language models store all knowledge in their parameters, leading to several fundamental limitations:

Consider the query: "What were the recent developments in quantum computing at IBM in 2024?"

Parametric LLMs face the following critical challenges:

- **Knowledge Cutoff:** Training data has temporal boundaries
- **Hallucination:** Models generate plausible but incorrect information
- **Parameter Inefficiency:** Storing all world knowledge in weights is computationally prohibitive
- **Update Difficulty:** Incorporating new information requires expensive retraining

Why RAG ?

RAG solves these limitations by augmenting generation with real-time retrieval from external knowledge bases, combining the fluency of large language models with the accuracy of information retrieval systems.

RAG Architecture: Mathematical Framework

RAG operates by retrieving relevant documents from an external corpus and conditioning the language model’s generation on this retrieved context.

Core Mathematical Formulation:

Let x be the input query, $D = \{d_1, d_2, \dots, d_N\}$ be the external knowledge corpus, and y be the generated response. RAG computes:

$$P_{RAG}(y|x) = \sum_{z \subset D} P(z|x) \cdot P(y|x, z) \quad (1)$$

Where:

- $P(z|x)$ is the retrieval distribution over document subsets z
- $P(y|x, z)$ is the generation probability conditioned on query x and retrieved documents z

Practical Implementation:

$$\text{Retrieved Documents: } z = \text{top-k}(\text{similarity}(x, D)) \quad (2)$$

$$\text{Generation: } y \sim P_{LM}(y|x \oplus z) \quad (3)$$

Where $\oplus$ denotes context concatenation and P_{LM} is the language model distribution.

Two-Stage Process: RAG fundamentally decomposes the generation problem into two sequential stages:

- **Retrieval Stage** - finding relevant documents using similarity search,
- **Generation Stage** - producing responses conditioned on retrieved context.

This decomposition allows independent optimization of each component while maintaining end-to-end differentiability.

Dense Retrieval: Vector Space Mathematics

RAG relies on dense retrieval using learned embeddings to find semantically relevant documents.

Embedding Generation:

$$\mathbf{q} = \text{Encoder}_Q(x) \in \mathbb{R}^d \quad (\text{Query embedding}) \quad (4)$$

$$\mathbf{d}_i = \text{Encoder}_D(d_i) \in \mathbb{R}^d \quad (\text{Document embedding}) \quad (5)$$

Similarity Computation:

$$\text{sim}(x, d_i) = \frac{\mathbf{q} \cdot \mathbf{d}_i}{\|\mathbf{q}\| \|\mathbf{d}_i\|} = \cos(\mathbf{q}, \mathbf{d}_i) \quad (6)$$

Retrieval Probability:

$$P(d_i|x) = \frac{\exp(\text{sim}(x, d_i)/\tau)}{\sum_{j=1}^N \exp(\text{sim}(x, d_j)/\tau)} \quad (7)$$

Where τ is a temperature parameter controlling the sharpness of the distribution.

Alternative Similarity Metrics: Beyond cosine similarity, RAG systems may employ:

- **Dot Product:** $\text{sim}(x, d_i) = \mathbf{q} \cdot \mathbf{d}_i$ (computationally efficient)
- **Euclidean Distance:** $\text{sim}(x, d_i) = -\|\mathbf{q} - \mathbf{d}_i\|_2$ (captures magnitude differences)
- **Learned Similarity:** $\text{sim}(x, d_i) = f_\theta(\mathbf{q}, \mathbf{d}_i)$ where f_θ is a neural network

Embedding Space Properties: The quality of retrieval depends critically on the embedding space geometry. Effective embeddings satisfy:

- **Semantic Clustering** - semantically similar documents form dense clusters
- **Query-Document Alignment** - queries and relevant documents have high cosine similarity
- **Negative Separation** - irrelevant documents are well-separated from queries

RAG Pipeline Architecture

The complete RAG system architecture is illustrated in Figure 1, showing the flow from user query to final response generation.




Figure 1: RAG pipeline showing query encoding, retrieval, and augmented generation

The architecture demonstrates two key phases: **Offline Indexing** (dashed arrow) where documents are encoded and stored in the vector database, and **Online Retrieval** where user queries trigger the retrieval-generation pipeline.

Worked Example: RAG for Scientific Q&A

Let's implement RAG for the query "What is the applications of transformer architecture?" with a corpus of 3 documents.

Setup:

- Embedding dimension: $d = 4$
- Number of documents: $N = 3$
- Top-k retrieval: $k = 2$
- Temperature: $\tau = 0.1$

Document Corpus:

$$d_1 : \text{"Transformers excel in natural language processing tasks"} \quad (8)$$

$$d_2 : \text{"Computer vision applications use transformer-based models"} \quad (9)$$

$$d_3 : \text{"Quantum computing requires specialized algorithms"} \quad (10)$$

Step 1: Generate Embeddings

Query embedding (simplified):

$$\mathbf{q} = \text{Encoder}_Q(\text{"transformer applications"}) = \begin{bmatrix} 0.8 \\ 0.6 \\ 0.2 \\ 0.4 \end{bmatrix} \quad (11)$$

Document embeddings:

$$\mathbf{d}_1 = \begin{bmatrix} 0.9 \\ 0.5 \\ 0.3 \\ 0.2 \end{bmatrix} \quad (\text{NLP focus}) \quad (12)$$

$$\mathbf{d}_2 = \begin{bmatrix} 0.7 \\ 0.7 \\ 0.1 \\ 0.6 \end{bmatrix} \quad (\text{Vision focus}) \quad (13)$$

$$\mathbf{d}_3 = \begin{bmatrix} 0.2 \\ 0.1 \\ 0.9 \\ 0.8 \end{bmatrix} \quad (\text{Quantum focus}) \quad (14)$$

Step 2: Compute Similarities

Cosine similarities:

$$\text{sim}(\mathbf{q}, \mathbf{d}_1) = \frac{\mathbf{q} \cdot \mathbf{d}_1}{\|\mathbf{q}\| \|\mathbf{d}_1\|} = \frac{0.8 \times 0.9 + 0.6 \times 0.5 + 0.2 \times 0.3 + 0.4 \times 0.2}{\sqrt{1.4} \times \sqrt{1.23}} = \frac{1.16}{1.31} = 0.885 \quad (15)$$

$$\text{sim}(\mathbf{q}, \mathbf{d}_2) = \frac{0.8 \times 0.7 + 0.6 \times 0.7 + 0.2 \times 0.1 + 0.4 \times 0.6}{\sqrt{1.4} \times \sqrt{1.35}} = \frac{1.42}{1.37} = 1.036 \quad (16)$$

$$\text{sim}(\mathbf{q}, \mathbf{d}_3) = \frac{0.8 \times 0.2 + 0.6 \times 0.1 + 0.2 \times 0.9 + 0.4 \times 0.8}{\sqrt{1.4} \times \sqrt{1.49}} = \frac{0.72}{1.44} = 0.500 \quad (17)$$

Step 3: Compute Retrieval Probabilities

$$P(d_1|x) = \frac{\exp(0.885/0.1)}{\exp(0.885/0.1) + \exp(1.036/0.1) + \exp(0.500/0.1)} = \frac{\exp(8.85)}{\exp(8.85) + \exp(10.36) + \exp(5.00)} \quad (18)$$

$$= \frac{6946}{6946 + 31855 + 148} = 0.178 \quad (19)$$

$$P(d_2|x) = \frac{\exp(10.36)}{38949} = 0.818 \quad (20)$$

$$P(d_3|x) = \frac{\exp(5.00)}{38949} = 0.004 \quad (21)$$

Step 4: Select Top-k Documents

Based on similarities: $\text{sim}(\mathbf{q}, \mathbf{d}_2) > \text{sim}(\mathbf{q}, \mathbf{d}_1) > \text{sim}(\mathbf{q}, \mathbf{d}_3)$

Retrieved documents: $z = \{d_2, d_1\}$

Step 5: Form Augmented Context

$$c = x \oplus z = \text{"Query: transformer applications \n Context: "} + d_2 + d_1 \quad (22)$$

Step 6: Generate Response

The language model generates:

$$y \sim P_{LM}(y|c) = \text{"Transformers have applications in NLP and computer vision..."} \quad (23)$$

Mathematical Insight: The temperature parameter $\tau = 0.1$ creates a sharp distribution, heavily favoring d_2 with 81.8% probability. Higher temperatures (e.g., $\tau = 1.0$) would create more uniform distributions, potentially including lower-similarity documents. This demonstrates the critical role of hyperparameter tuning in RAG systems.

Summary

RAG revolutionizes language models by mathematically combining parametric knowledge with non-parametric retrieval:

- Enables dynamic knowledge access without parameter updates
- Reduces hallucination through evidence-based generation
- Provides interpretable attribution via retrieved documents
- Maintains computational efficiency through external knowledge storage

The mathematical elegance lies in the probabilistic framework that seamlessly integrates retrieval distributions with generation probabilities. This approach addresses fundamental limitations of parametric models while preserving their generative capabilities, establishing RAG as a cornerstone technique for knowledge-intensive AI applications.

Practical Implementation Notes: Real-world RAG systems must address several mathematical challenges:

- **Retrieval Latency** - balancing accuracy with speed through approximate search algorithms
- **Domain Adaptation** - fine-tuning embeddings for specialized corpora
- **Multi-hop Reasoning** - extending RAG to iterative retrieval for complex queries requiring multiple information sources